

# Deqing Fu

**Email:** deqingfu@usc.edu    **Address:** Ginsburg Hall, USC, Los Angeles, CA 90089  
**Website:** deqingfu.github.io    **Office:** GCS 403E

<b>Education</b>	<b>University of Southern California</b> <i>Ph.D. in Computer Science</i> Advisors: Vatsal Sharan & Robin Jia	2022 - Present
	<b>University of Chicago</b> <i>S.M. in Statistics</i>	2020 - 2022
	<b>University of Chicago</b> <i>S.B. in Mathematics</i> with Honors <i>S.B. in Computer Science</i> with Honors <i>A.B. in Statistics</i>	2016 - 2020
<b>Experience</b>	<b>Bluelight AI</b> <i>Student Researcher</i> Hosts: John Carlsson, Jakob Hansen, Gunnar Carlsson	Jun. 2025 – Present
	<b>Simons Institute at UC Berkeley</b> <i>Visiting Graduate Student</i> Program: Modern Paradigms in Generalization	Fall 2024
	<b>Meta Llama Team</b> <i>Research Scientist Intern</i> → <i>Student Researcher</i> Host: Lawrence Chen; Peers: Tong Xiao, Rui Wang, and Pengchuan Zhang.	May – Dec., 2024
	<b>Google</b> <i>Software Engineering Intern</i> Teams: Image Search; Knowledge Engine; Bard (now Gemini) Multilingual.	Summer 2021, 2022, 2023
<b>Publications</b>  * denotes equal contribution	<b>TLDR: Token-Level Detective Reward Model for Large Vision Language Models</b> Deqing Fu, Tong Xiao, Rui Wang, Wang Zhu, Pengchuan Zhang, Guan Pang, Robin Jia, Lawrence Chen	ICLR 2025
	<b>Transformers Learn Low Sensitivity Functions: Investigations and Implications</b> Bhavya Vasudeva*, Deqing Fu*, Tianyi Zhou, Elliot Kau, You-Qi Huang, Vatsal Sharan	ICLR 2025
	<b>DeLLMa: Decision Making Under Uncertainty with Large Language Models</b>	

Ollie Liu\*, Deqing Fu\*, Dani Yogatama, Willie Neiswanger  
*ICLR 2025 Spotlight (Top 5.1%)* ICLR 2025

**DreamSync: Aligning Text-to-Image Generation with Image Understanding Feedback**

Jiao Sun\*, Deqing Fu\*, Yushi Hu\*, Su Wang, Royi Rassin, Da-Cheng Juan, Dana Alon, Charles Herrmann, Sjoerd van Steenkiste, Ranjay Krishna, Cyrus Rashtchian NAACL 2025

**Transformers Learn to Achieve Second-Order Convergence Rates for In-Context Linear Regression**

Deqing Fu, Tian-Qi Chen, Robin Jia, Vatsal Sharan  
*SoCalNLP Symposium 2023 Best Paper Award* NeurIPS 2024

**Pre-trained Large Language Models Use Fourier Features to Compute Addition**

Tianyi Zhou, Deqing Fu, Vatsal Sharan, Robin Jia NeurIPS 2024

**IsoBench: Benchmarking Foundation Models on Isomorphic Representations**

Deqing Fu\*, Ruohao Guo\*, Ghazal Khalighinejad\*, Ollie Liu\*, Bhuwan Dhingra, Dani Yogatama, Robin Jia, Willie Neiswanger  
*Huggingface 🗝️: [ Dataset ]* COLM 2024

**SCENE: Self-Labeled Counterfactuals for Extrapolating to Negative Examples**

Deqing Fu, Ameya Godbole, Robin Jia EMNLP 2023

**Harnessing the Conditioning Sensorium for Improved Image Translation**

Cooper Nederhood, Nicholas Kolkin, Deqing Fu, Jason Salavon ICCV 2021

**Preprints**

**Resa: Transparent Reasoning Models via SAEs**

Shangshang Wang, Julian Asilis, Ömer Faruk Akgül, Enes Burak Bilgin, Ollie Liu, Deqing Fu, Willie Neiswanger

**Zebra-CoT: A Dataset for Interleaved Vision Language Reasoning**

Ang Li\*, Charles L. Wang\*, Kaiyu Yue\*, Zikui Cai\*, Ollie Liu\*, Deqing Fu\*, Peng Guo\*, Wang Bill Zhu\*, Vatsal Sharan, Robin Jia, Willie Neiswanger, Furong Huang, Tom Goldstein, Micah Goldblum  
*Huggingface 🗝️: [ Dataset ] [ Models ]*

**Textual Steering Vectors Can Improve Visual Understanding in Multimodal Large Language Models**

Woody Haosheng Gan\*, Deqing Fu\*, Julian Asilis\*, Ollie Liu\*,  
Dani Yogatama, Vatsal Sharan, Robin Jia, Willie Neiswanger

**FoNE: Precise Single-Token Number Embeddings via  
Fourier Features**

Tianyi Zhou, Deqing Fu, Mahdi Soltanolkotabi, Robin Jia, Vatsal Sharan

**VisualLens: Personalization through Visual History**

Wang Bill Zhu, Deqing Fu, Kai Sun, Yi Lu, Zhaojiang Lin, Seungwhan Moon, Kanika Narang, Mustafa Canim, Yue Liu, Anuj Kumar, Xin Luna Dong

**Awards**

<b>Viterbi Best Research Assistant Award</b> ( <i>one per department</i> )	USC, 2025
<b>SoCalNLP Symposium 2023 Best Paper Award</b>	SoCalNLP, 2023
<b>Provost's Fellowship</b>	USC, 2022
<b>Susanne H. Rudolph Scholarship</b>	UChicago, 2020-2021
<b>Liew Family College Research Fellows Fund</b>	UChicago, 2020
<b>Dean's List</b>	UChicago, 2016-2020

**Grants**

**VESSL AI Academia Program**, 1,200 AI100 Hours ( $\approx$  \$5,000 value)  
PI: Deqing Fu August 2025  
Title: *Emergence of Fourier Features in Large Language Model Pretraining*

**NSF ACCESS Allocation**, 750,000 Credits ( $\approx$  \$75,000 value)  
PI: Deqing Fu, Co-PI: Robin Jia August 2025 - August 2026  
Title: *Robust Mathematical Reasoning with Fourier Language Models*

**Invited Talks**

<i>Algorithmic Perspectives on Understanding Transformers</i>	
FlaNN Talk Seminar	Aug 25, 2025
<i>Closing the Modality Gap:</i>	
<i>Benchmarking and Improving Visual Understanding in Multimodal LLMs</i>	
Stanford NLP Seminar	May 22, 2025
<i>Algorithmic Perspectives on Understanding Transformers</i>	
UChicago/TTIC NLP Seminar	May 08, 2025
<i>Tutorial on the interplay between Theory and Interpretability</i>	
USC ML/NLP Joint Seminar	Apr 25, 2025
<i>Algorithmic Perspectives on Understanding Transformers</i>	
Duke NLP Seminar	Feb 12, 2025
<i>Guest Lecture on Interpretability and AI Safety</i>	
DSO 575: Driving Business Transformation with GenAI	Oct 24, 2024

	<i>Simplicity Bias or Transformers to Learn Low Sensitivity Functions</i> USC NLP Lunch	Feb 22, 2024
	<i>Aligning Text-to-Image Generation with Image Understanding Feedback</i> USC NLP Lunch	Oct 19, 2023
	<i>Transformers Learn to Achieve Second-Order Convergence for In-Context Linear Regression</i> USC NLP Lunch	Mar 29, 2023
<b>Services</b>	<b>Reviewer</b> for conferences in <ul style="list-style-type: none"> <li>• <b>Machine Learning:</b> ICLR 2025, NeurIPS (2024, 2025), ICML 2025;</li> <li>• <b>Natural Language Processing:</b> COLM 2024, EMNLP (2023, 2024), ACL (2024, 2025);</li> <li>• <b>Computer Vision:</b> CVPR (2023, 2024, 2025), ECCV 2022, ICCV 2025</li> </ul>	
	<b>Organizer</b> of USC ML/NLP Joint Reading Group	Jan 2025 – Present
	<b>Organizer</b> of USC AI Safety Club	Jan 2024 – Jan 2025
<b>Students Supervised</b>	<b>Master's and Undergraduate Students</b>	
	Gustavo Carvalho (USC MS)	Feb 2024 – Present
	Qilin Ye (USC UG $\rightarrow$ Duke MS)	Sep 2023 – Present
	Woody Haosheng Gan (USC UG)	Sep 2023 – Present
	You-Qi Huang (USC UG)	Sep 2023 – Mar 2024
	Elliot Kau (USC MS)	Sep 2023 – Mar 2024
	Tian-Qi Chen (USC MS)	Mar 2023 – Sep 2023
<b>Teaching Experience</b>	<b>Teaching Assistant</b> , CSCI 467: Introduction to Machine Learning University of Southern California, Los Angeles, CA	
		Spring 2025
	<b>Teaching Assistant</b> , CSCI 566: Deep Learning and its Applications University of Southern California, Los Angeles, CA	
		Spring 2023
	<b>Teaching Assistant</b> , BUSN 41301: Statistical Insight Booth School of Business, University of Chicago, Chicago, IL	
		Fall 2021
	<b>Course Assistant</b> , CMSC 31230: Fundamentals of Deep Learning University of Chicago, Chicago, IL	
		Fall 2021
	<b>Course Assistant</b> , MATH 20300: Analysis in $\mathbb{R}^n$ - I University of Chicago, Chicago, IL	
		Spring 2018
<b>Skills</b>	C/C++, Python, $\text{\LaTeX}$ , MATLAB, R, Java, CUDA PyTorch, TensorFlow, Keras, JAX, HuggingFace	